

Snapdragon Local Review

- Private, offline-first code and Vercel preflight
- NPU-accelerated developer feedback for Snapdragon PCs

01 | SNAPDRAGON LOCAL REVIEW

The problem

- Cloud code review creates latency and data-governance risk
- Manual pre-deployment checks miss secrets and configuration errors
- Local builds already consume the CPU developers need for feedback

02 | SNAPDRAGON LOCAL REVIEW

The solution

- A single local CLI before commit or deployment
- Deterministic checks catch actionable issues immediately
- An optional local coding model explains and remediates findings
- No API keys, no source upload, no cloud fallback

03 | SNAPDRAGON LOCAL REVIEW

Technical architecture

- Repository → CLI → deterministic rules
- Optional adapter → tokenizer → ONNX Runtime
- QNN Execution Provider → Hexagon NPU / HTP
- CPU remains available for compilers, Docker, tests, and local servers

Why Snapdragon

- Privacy: proprietary code remains on-device
- Responsiveness: inference is offloaded from the CPU
- Offline operation: useful on flights and restricted networks
- Enterprise fit: local governance and predictable data boundaries

Developer workflow

- snapdragon-review check .
- JSON output integrates with editors and CI
- --strict turns warnings into deployment blockers
- Pre-commit catches issues before staging or Vercel

Demo and measurement

- Inject a token and insecure Vercel rewrite
- Run the CLI and show structured findings
- Show QNN provider availability on ARM64

- Compare review latency and CPU utilization during a build

07 | SNAPDRAGON LOCAL REVIEW

Accessibility and roadmap

- Works today without AI dependencies
- Add Transformers CPU fallback for compatibility

- Add ONNX Runtime GenAI/QNN adapter for NPU inference
- Expand rules, VS Code integration, and signed model bundles

08 | SNAPDRAGON LOCAL REVIEW

Call to action

- Install locally

- Keep the review loop private
- Turn every Snapdragon developer machine into an edge code-review station